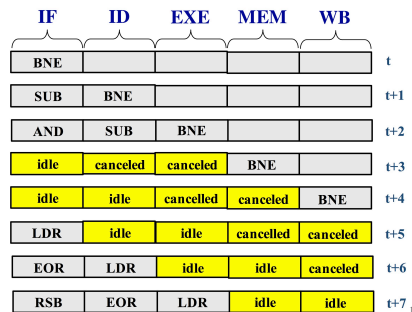


## Pipelined Instruction Execution

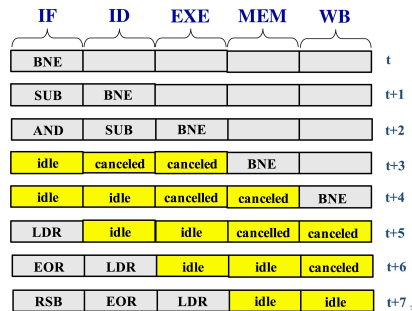


## In this Lesson

- Principles of pipelined instruction execution
- Pipeline hazards
  - Data forwarding
  - Loop Buffer
  - Delayed branches
  - Branch prediction
  - Multiple prefetching
- Performance with pipelined instruction execution units

2

## Pipelined Instruction Execution Part I: Pipelines and Hazards



3

## Execution Without Pipeline Stages

- The execution unit takes one instruction at a time.
- If an instruction takes  $n$  cycles to execute,  $m$  instructions take  $n \times m$  cycles to complete.

4

## Pipelining

- Pipelining is executing instructions by stages like a production line



- Ideally, the instructions pass from one stage to the other on every cycle and are executed at a rate of one per cycle ( $m$  instructions take  $m$  cycles to complete)

5

## Pipeline of Five Stages

- **Instruction Fetch (IF)**
  - Brings the instruction from memory into the CPU.
- **Instruction decode/register fetch (ID)**
  - Decodes the instruction,
  - Reads the content of the registers specified as operands into temporary registers
  - Sign-extends the immediate values,
  - Determines the target address of branch instructions
- **Execution/effective address (EX)**
  - Load/stores – calculates the effective address
  - Arithmetic/logic – executes the operation

6

## Pipeline of Five Stages

### • Memory access (MEM)

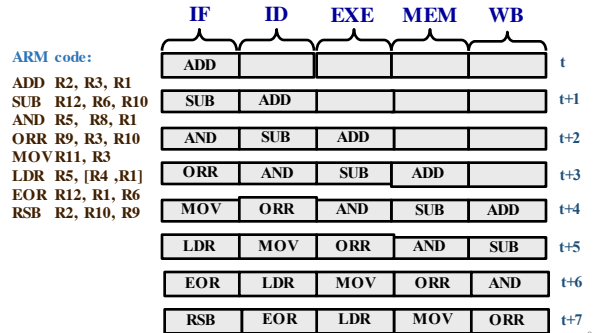
- Load – reads the effective address
- Store – writes the effective address
- Any other instruction – no action

### • Write-back (WB)

- Writes the result in the destination register if it is an arithmetic/logic or a load instruction

7

## Ideal Execution of Instructions on a Pipelined Unit



## Pipeline Hazards (Interlocks)

- Are the main reason why pipelines do not exhibit ideal performance.
- Are basically instructions interdependencies that have the potential to produce an incorrect outcome of a program.
- May cause instructions to stall in the pipelined until the hazard is resolved
- Have a negative effect on the performance of the processor due to the wasted idle cycles caused by the stalling instructions.

9

## Types of Hazards

- Structural hazards
- Data Hazards
- Control Hazards

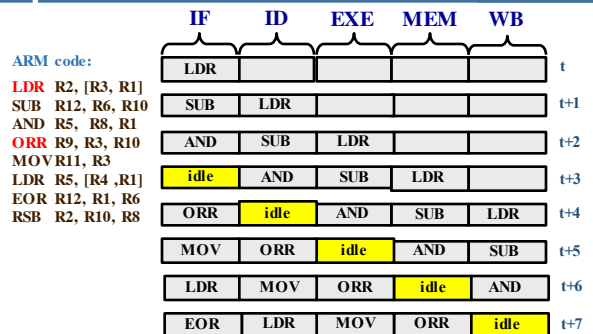
10

## Structural Hazards

- Caused by more than one instructions attempting to access the same resource (memory for example) at the same time.
- In a single processor system it happens when a load/store instruction is in the **MEM** stage and the processors attempts to fetch another instruction in the **IF** stage.
- In a multiprocessor system the hazard could arise from two or more processors trying to access memory at the same time.
- To prevent this hazard the load/store instruction in the **MEM** stage is allowed to proceed and the instruction that was to be fetched is prevented from entering the **IF** stage.
- A structural hazard is present any time a load/store instruction is in the **MEM** stage.

11

## Example of Structural Hazard for Single Port Memory



## Solutions for Structural Hazards

- Using memory systems with more than one memory port (at least one for instructions and one for data).
- Using split caches (one cache for instructions and another for data)
- Both eliminate the structural hazard penalty of one cycle.

13

## Data Hazards

- Occur when data reads and writes could take place in an order different from the order prescribed by the program
- This could happen due to the structure of the pipeline stages or in processors that allows out-of order instruction execution
- There are three types of data hazards:
  - Read after Write (RAW)
  - Write after Read (WAR)
  - Write after Write (WAW)

14

## RAW Hazard

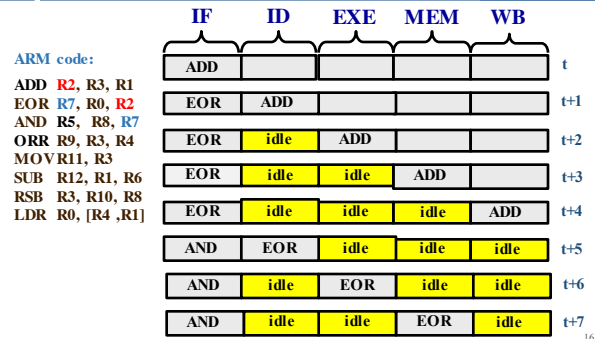
- Arises when an instruction attempts to read the content of a source register that is the destination register of a previous instruction that has not completed the writing of its destination register.

```
ADD R2, R3, R1
EOR R7, R9, R2
```

- In a pipeline of five stages this hazard is detected when the EOR instruction attempts to enter the ID stage while the ADD instruction has not left the WB stage.
- A simple way to prevent this hazard is to halt the EOR instruction in the IF stage until the ADD instruction leaves the WB stage.

15

## Example of RAW Hazards



## WAR Hazard

- Arises when an instruction attempts to write a register that is a source register of a previous instruction that has not completed the reading of its source registers.

```
ADD R2, R3, R7
EOR R7, R9, R2
```

- Can only happen in processors that allow out-of-order instruction execution.
- Not the case in ARM because the reading of registers always take place before the writing of a register of an instruction that follows.

17

## WAW Hazard

- Arises when an instruction attempts to write a register that is a destination register of a previous instruction before that instruction completes the writing of its destination registers.

```
ADD R7, R3, R7
EOR R7, R9, R2
```

- Can only happen in processors that allow out-of-order instruction execution (not the case in ARM)

18

## Solution for Data Hazards: Instructions Reordering

ARM code:

```
ADD R2, R3, R1
EOR R7, R0, R2
AND R5, R8, R7
ORR R9, R3, R4
MOVR11, R3
SUB R12, R1, R6
RSB R3, R10, R8
LDR R0, [R4, R1]
```

ARM code:

```
ADD R2, R3, R1
ORR R9, R3, R4
MOVR11, R3
SUB R12, R1, R6
EOR R7, R0, R2
AND R5, R8, R7
RSB R3, R10, R8
LDR R0, [R4, R1]
```

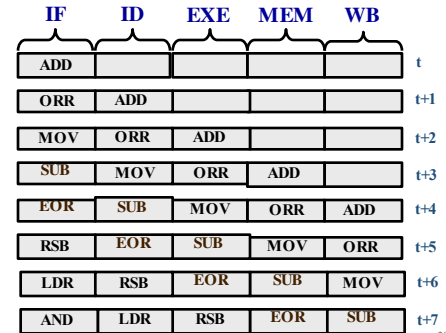
ARM code:

```
ADD R2, R3, R1
ORR R9, R3, R4
MOVR11, R3
SUB R12, R1, R6
EOR R7, R0, R2
RSB R3, R10, R8
LDR R0, [R4, R1]
AND R5, R8, R7
```

## Solution for Data Hazards: Instructions Reordering

ARM code:

```
ADD R2, R3, R1
ORR R9, R3, R4
MOVR11, R3
SUB R12, R1, R6
EOR R7, R0, R2
RSB R3, R10, R8
LDR R0, [R4, R1]
AND R5, R8, R7
```



## Solution for Data Hazards: Data Forwarding

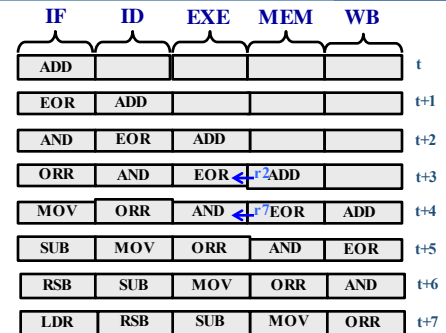
- Allows the pipeline to use the value of a destination register after it is generated (without having to wait until after the WB stage)
  - The value of a destination register is generated in the EXE stage for arithmetic/logic instructions.
  - The value of a destination register is generated in the MEM stage for load/store instructions.
- Eliminates all the penalty cycles due to the data hazard prevention
- Allows the dependent instruction to advance to the ID stage even if a source registers has not been updated

21

## Example of Data Forwarding

ARM code:

```
ADD R2, R3, R1
EOR R7, R0, R2
AND R5, R8, R7
ORR R9, R3, R4
MOVR11, R3
SUB R12, R1, R6
RSB R3, R10, R8
LDR R0, [R4, R1]
```

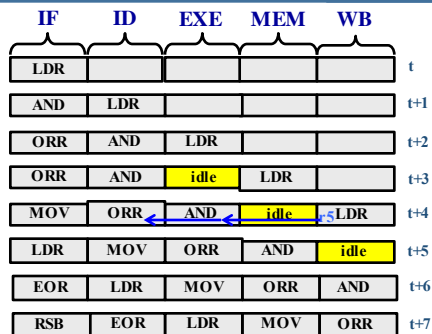


22

## Another Example of Data Forwarding

ARM code:

```
LDR R5, [R3, R1]
AND R5, R8, R5
ORR R9, R3, R5
MOVR11, R3
LDR R5, [R4, R1]
EOR R12, R1, R6
RSB R2, R10, R8
```



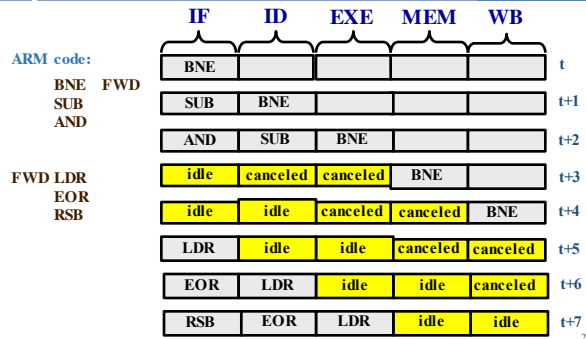
23

## Control Hazards

- Typically, the instructions following a branch instruction in a program are fed into the pipeline following the branch instruction.
- If the branch is not taken there is no penalty
- If the branch is taken, the instructions following the branch on the pipeline must be cancelled.
- After the instructions are cancelled the branch target instructions must be supplied to the pipeline
- The cycles wasted as a result of the cancellation of the instructions plus the cycles it takes for the target instructions to start entering the pipeline after the cancellation of the instructions (the memory latency) constitute the branch execution penalty.

24

## Control Hazard Example (memory latency of two cycles)



## Alternatives for Reducing Branch Execution Time

- Loop buffer
- Delayed branch
- Branch prediction
- Multiple Prefetching
- Combinations of the above

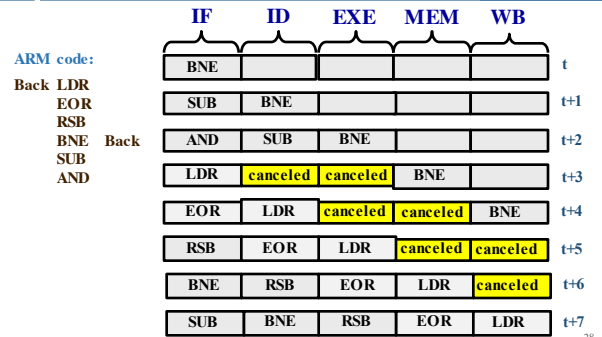
26

## Loop Buffer

- A small, very fast cache is used to keep the previous instructions.
- The objective of this cache is to hold the instructions corresponding to a loop so that the target address can be fetched immediately after the branch condition is tested.
- It reduces the memory latency for fetching the target instructions.
- Still pays a penalty for the instructions following the branch that must be canceled

27

## Effect of Loop Buffer (zero memory latency)

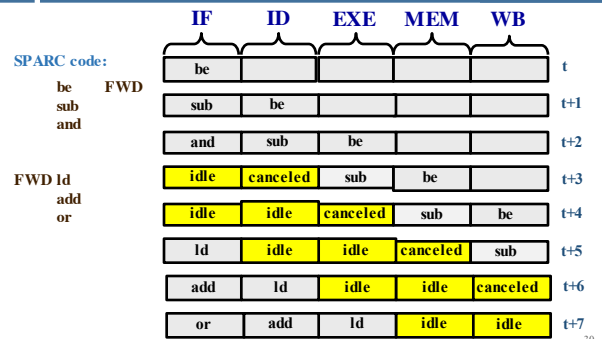


## Delayed Branch

- If the condition of the branch is asserted a branch to the target address is taken after the instruction following the branch is executed.
- The machine code programmer or the compiler will try to place a useful instruction after the branch instruction.
- If an useful instruction cannot be placed after the branch instruction a no op instruction must be placed instead (There will be no benefit).
- SPARC and MIPS support this type of branch

29

## Delayed Branch Example (memory latency of two cycles)



## Effectiveness of Delayed Branch

- Depends on the compiler
  - Place *useful instructions* after the branch instructions
  - Place *nop instructions* after the branch instructions
- Improves with delay of more than one instruction

31

## Branch Prediction

- The instruction fetch unit always make a prediction of what will be the result of the evaluation of the condition of a conditional branch.
- The instructions that are sent to the pipeline following the branch depend on the prediction

### Prediction that the branch will take place

The stream of instructions at the target address are sent immediately after the branch instruction.

### Prediction that the branch will no take place

The stream of instructions following the branch are sent immediately after the branch instruction.

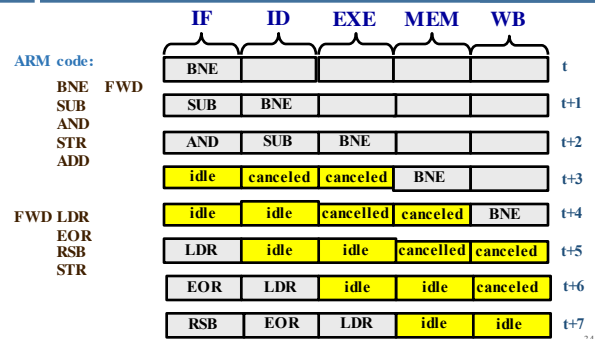
32

## Outcome of Branch Prediction

- If the prediction is asserted there is no branch penalty
- If the prediction is not asserted:
  - The instructions sent to the pipeline after the branch are canceled.
  - The other stream of instructions is sent to the pipeline with a delay caused by the memory latency (the time it takes memory to respond).

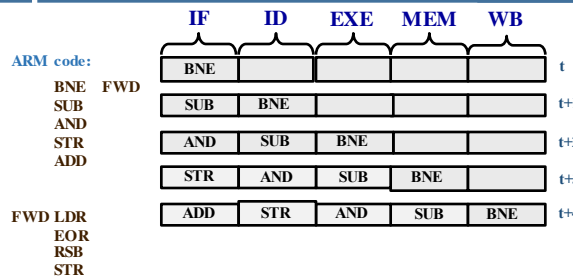
33

## Prediction: Branch will not take place Outcome of Prediction: Not Asserted



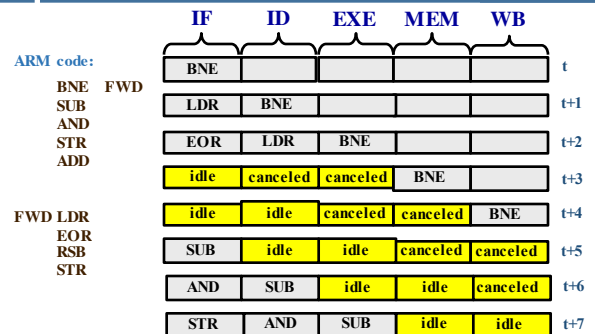
34

## Prediction: Branch will not take place Outcome of Prediction: Asserted



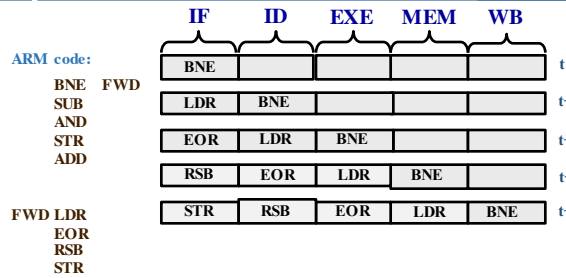
35

## Prediction: Branch will take place Outcome of Prediction: Not Asserted



36

Prediction: **Branch will take place**  
Outcome of Prediction: **Asserted**



37

## Types of Branch Predictions

- Static Prediction
  - The prediction is specified by the branch instruction
  - The compiler decides the prediction
- Dynamic Prediction
  - The hardware determines the prediction base on previous branch outcomes (taken/not taken)

38

## Static Prediction Schemes

- Predict backward branches taken and forward branches not taken
- Hint Bits
  - Some architectures allow the compiler to specify some bits for a branch instruction that hints the outcome of the evaluation of its condition

39

## Dynamic Prediction Schemes

- N-bit Saturated Counter Predictors
  - 1-bit History Predictor
  - Two-bit Saturated Counter Predictor
- Two Level Branch Predictor

40


## N-bit Saturated Counter Predictors

- An n-bit counter is incremented if the outcome of a branch is taken and decremented if the the outcome is not taken.
- When the counter reaches the maximum count and the next branch outcome is taken the count remains the maximum (it saturates).
- When the counter reaches the minimum count and the next branch outcome is not taken the count remains the minimum (it saturates).
- The prediction is not taken if the count is less than  $2^{n-1}$ . Otherwise the prediction is taken.

41

## 1 bit History Predictor

- Predicts the same branch outcome the last time it was executed
  - Set to one if a branch is taken (a one predicts taken)
  - Set to zero if a branch is not taken (a zero predicts not taken)




42

## Two-bit Saturated Counter Predictor

- Predicts with hysteresis (takes into account the history of a branch outcome further than the last time it was executed)
  - Counter incremented if branch taken (no further incrementation when a count of three is reached)
  - Counter decremented if branch not taken (no further decrementation when a count of zero is reached)
  - If the most significant bit of the counter is one predict branch taken, otherwise predict branch not taken

43

## Two-bit Saturated Counter Predictor



44

## Issues with N-bit Saturated Counter Predictors

- The saturated counters are kept on a one-dimensional table of  $2^k$  entries.
- The branches are mapped to the table using the least significant  $k$  bits of the branch address and one of the following types of mapping:
  - Direct mapping
  - Fully associative mapping
  - Set associative mapping
  - Hashing function
- Are prone to an aliasing problem because more than one branch could use the same counter (there is not one counter per branch address)

45

## Two Level Branch Predictor

- Takes into consideration that the outcome of a branch depends on the history of other recent branches and on the history of the branch itself
- Uses a shift register to shift in the results of branch outcomes (1 for taken and 0 for not taken)
- Uses the shift register to address a two-dimensional table of saturated counters.
- It is prone to an aliasing problem

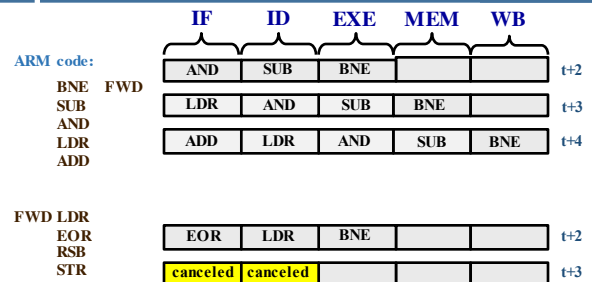
46

## Multiple Prefetching

- The instruction stream following the branch instruction and the branch target stream are executed in parallel pipelines until the branch condition is evaluated
- Once the condition is evaluated one of the instructions stream is cancelled and the other proceeds
- Requires a branch management unit

47

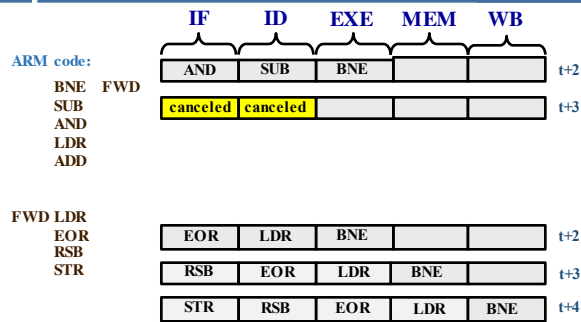
## Multiple prefetching: Branch not taken



48



## Multiple prefetching: Branch taken



49

## Disadvantages of Multiple Prefetching

- Introduces bus and registers contention
- May end up requiring additional pipelines when a branch is immediately followed by another branch

50

## Other Mechanisms for Reducing the Effect of Interlocks

- Superscalar
  - Rely on multiple function units that operate in parallel (adders, multipliers, dividers)
  - Potentially, can execute instructions at a rate higher than one per cycle
  - Instructions requiring more than one cycle for executing do not obstruct others from executing
- Superpipeline
  - Many pipeline stages that can allow to reduce the system clock period
  - May increase the penalty for branches.

51