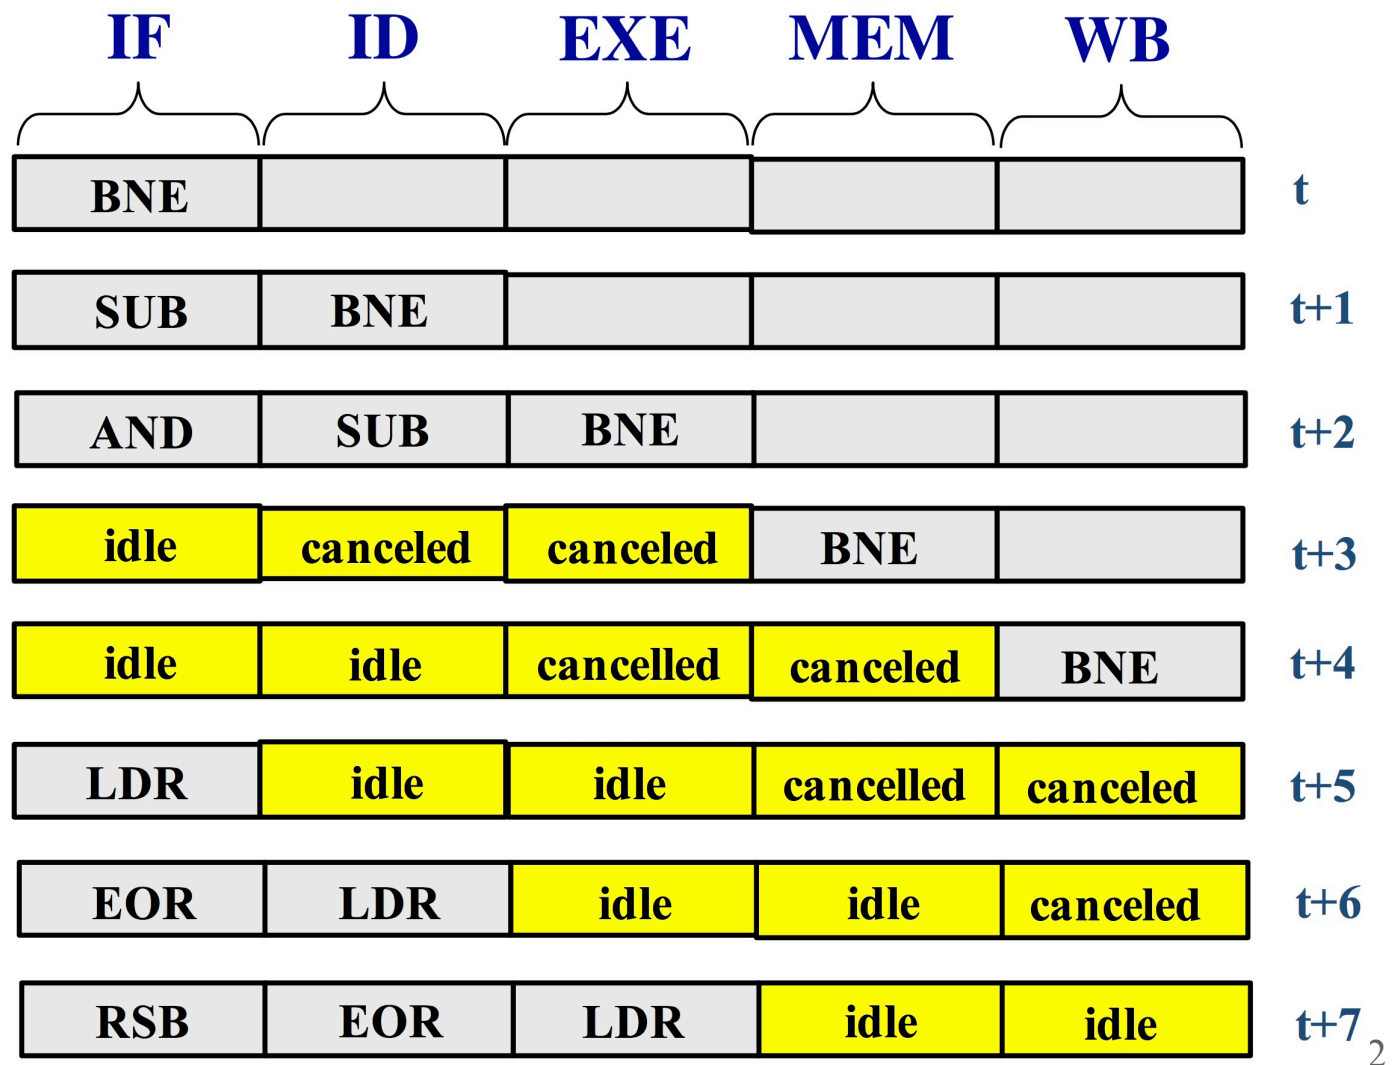


# Pipelined Instruction Execution

## Part II: Performance



# Performance Analysis

---

The number of cycles per instructions (CPI) is a standard measure of performance of microprocessors.

On an ideal pipelined unit, instructions execute at a rate of one per cycle. Thus, an ideal CPI is 1.

A simple way of calculating CPI is by determining the total number of cycles that take N instructions to execute (TC) and divide it by N.

Then,  $CPI = TC/N$

# Alternative Performance Analysis

The CPI can also be calculated in the following terms:

$$\text{CPI} = \text{CPI}_{\text{IDEAL}} + \text{CPI}_{\text{PENALTY}} = 1 + \text{CPI}_{\text{PENALTY}}$$

where  $\text{CPI}_{\text{PENALTY}}$  is the contribution of the extra wasted cycles caused by pipeline interlocks.

In general, the  $\text{CPI}_{\text{PENALTY}}$  is an aggregate of the wasted cycles introduced by structural, data and control hazards. Thus,

$$\text{CPI}_{\text{PENALTY}} = \text{Structural}_{\text{PENALTY}} + \text{Data}_{\text{PENALTY}} + \text{Control}_{\text{PENALTY}}$$

# The Control Penalty

Every branch has the potential for generating wasted cycles due to miss predictions

We can define the control penalty as follows:

Let  $PB$  = % of instructions that are branches

$BWC$  = average number of wasted cycles introduced by branch instructions.

Then,

$$\text{Control}_{\text{PENALTY}} = PB \times BWC$$

$$\text{Typically } BWC = \text{Prediction}_{\text{PENALTY}}$$

$$\text{Prediction}_{\text{PENALTY}} = (\% \text{ predictions not asserted}) \times (\# \text{ of cycles wasted on a miss prediction})$$

# Example Problem 1

The following piece of ARM code is executed 1000 times in a five stages pipeline with a Loop Buffer:

```
Back ADD
      SUB
      BNE Back
      ORR
      AND
```

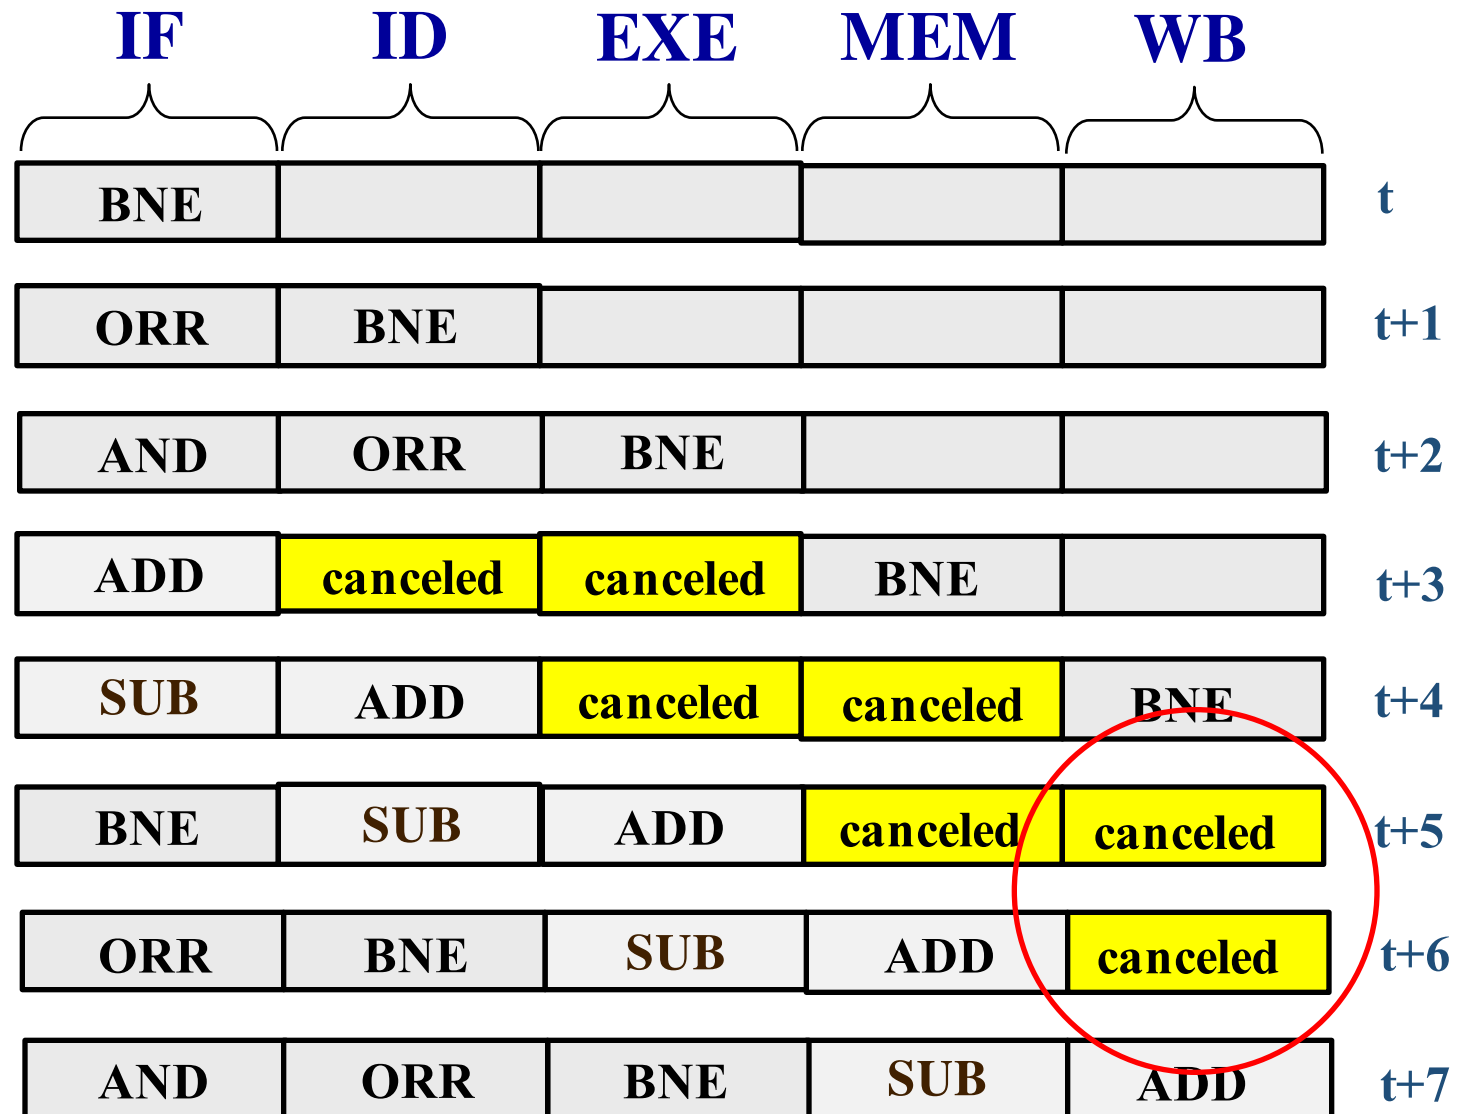
If there are no structural or data hazards, determine the CPI for each of the following branch prediction schemes:

- a) Normal ARM branch prediction (predict not taken)
- b) Static, backward branches taken/forward branches not taken
- c) 1 bit History Predictor
- d) Two-bit Saturated Counter Predictor

# Branch Penalty with a Loop Buffer

ARM code:

Back ADD  
SUB  
BNE Back  
ORR  
AND



# What we know

---

There are three instructions executed on each iteration of the loop, consuming one cycle each

$N = (\text{number of iterations}) \times (\text{numbers of instructions executed per iteration})$

$$N = 1000 \times 3 = 3000$$

There are 2 additional cycles wasted each time the prediction is not asserted

# Determining the Total Number of Wasted Cycles

Let

TC = total number of cycles that the 1000 iterations consume,

TIC = total number of cycles consumed by instructions

= (total numbers of iterations) x (number cycles consumed by instructions per iteration)

TWC = total number of wasted cycles

= (number of predictions not asserted) x (wasted cycles per prediction not asserted)

Then,  $TC = TIC + TWC$



# Normal ARM Branch

Since a branch not taken is predicted, the prediction is not asserted the first 999 iterations and asserted the last.

$TIC = (\text{total numbers of iterations}) \times (\text{number cycles consumed by instructions per iteration})$

$$TIC = 1000 \times 3 = 3000$$

$TWC = (\text{number of predictions not asserted}) \times (\text{wasted cycles per prediction not asserted})$

$$TWC = 999 \times 2 = 1998$$

$$TC = TIC + TWC = 3000 + 1998 = 4998$$

$$CPI = TC/N = 4998/(3000) = \mathbf{1.67}$$

# Static Prediction: Backward Branches Taken

---

Since a branch taken is predicted the prediction is asserted the first 999 iterations and not asserted the last.

$$\text{TIC} = 1000 \times 3 = 3000$$

$$\text{TWC} = 1 \times 2 = 2$$

$$\text{TC} = \text{TIC} + \text{TWC} = 3000 + 2 = 3002 \text{ cycles}$$

$$\text{CPI} = \text{TC}/\text{TI} = 3002 / 3000 = \mathbf{1.00067}$$

# 1 Bit History Predictor

The behavior of the one bit history counter will be as follows:

Iteration	Counter Before	Prediction	Outcome of Prediction	Counter After
1	0	Not taken	Not asserted	1
2	1	Taken	Asserted	1
3	1	Taken	Asserted	1
4	1	Taken	Asserted	1
5	1	Taken	Asserted	1
999	1	Taken	Asserted	1
1000	1	Taken	Not asserted	0

# 1 Bit History Predictor

---

$$\text{TIC} = 1000 \times 3 = 3000$$

$$\text{TWC} = 2 \times 2 = 4$$

$$\text{TC} = \text{TIC} + \text{TWC} = 3000 + 4 = 3004 \text{ cycles}$$

$$\text{CPI} = \text{TC}/\text{TI} = 3004 / 3000 = \mathbf{1.0013}$$

# Two-bit Saturated Counter Predictor

The behavior of the saturated counter will be as follows:

Iteration	Counter Before	Prediction	Outcome of Prediction	Counter After
1	00	Not taken	Not asserted	01
2	01	Not taken	Not asserted	10
3	10	Taken	Asserted	11
4	11	Taken	Asserted	11
5	11	Taken	Asserted	11
999	11	Taken	Asserted	11
1000	11	Taken	Not asserted	10

# Two-bit Saturated Counter Predictor

---

$$\text{TIC} = 1000 \times 3 = 3000$$

$$\text{TWC} = 3 \times 2 = 6$$

$$\text{TC} = \text{TIC} + \text{TWC} = 3000 + 6 = 3006 \text{ cycles}$$

$$\text{CPI} = \text{TC}/\text{TI} = 3006 / 3000 = \mathbf{1.002}$$

# Example Problem 1: Alternative Analysis

$$\text{CPI} = 1 + \text{CPI}_{\text{PENALTY}}$$

$$\text{CPI} = 1 + \text{Structural}_{\text{PENALTY}} + \text{Data}_{\text{PENALTY}} + \text{Control}_{\text{PENALTY}}$$

For this case there is no structural or data hazards. Then,

$$\text{CPI} = 1 + \text{Control}_{\text{PENALTY}} = 1 + \text{PB} \times \text{BWC}, \text{ where}$$

PB = % of instructions that are branches

BWC = average number of wasted cycles introduced by branch instructions

Since the program is a loop of three instructions PB = 33% = .33

BWC = (% predictions not asserted) x (# of cycles wasted on a miss prediction)

# Normal ARM Branch: Alternative Analysis

---

$BWC = (\% \text{ predictions not asserted}) \times (\# \text{ of cycles wasted on a miss prediction})$

$$BWC = (999/1000) \times 2 = 1.998$$

$$CPI = 1 + PB \times BWC = 1 + .33 \times 1.998 = 1.67$$



# Static Taken Prediction: Alternative Analysis

---

$BWC = (\% \text{ predictions not asserted}) \times (\# \text{ of cycles wasted on a miss prediction})$

$$BWC = (1/1000) \times 2 = .002$$

$$CPI = 1 + PB \times BWC = 1 + .33 \times .002 = 1.00067$$

# 1 bit History Predictor: Alternative Analysis

---

BWC = (% predictions not asserted) x (# of cycles wasted on a miss prediction)

$$\text{BWC} = (2/1000) \times 2 = .004$$

$$\text{CPI} = 1 + \text{PB} \times \text{BWC} = 1 + .33 \times .004 = 1.0013$$

# Two-bit Saturated Counter Predictor: Alternative Analysis

---

$BWC = (\% \text{ predictions not asserted}) \times (\# \text{ of cycles wasted on a miss prediction})$

$$BWC = (3/1000) \times 2 = .006$$

$$CPI = 1 + PB \times BWC = 1 + .33 \times .006 = 1.002$$

# Summary of Prediction Schemes

	Predictions Not Asserted	Total Number Waisted Cycles	CPI
ARM Branch	999	1998	1.67
Static Predict Taken	1	2	1.00067
1 Bit History Predictor	2	4	1.0013
Two-bit Saturated Counter	3	6	1.002

- Although for this example the Two-bit Saturated Counter predictor did not result in better performance than the 1 bit history predictor and the static predictor, it is superior to them because it can better adapt to different branching outcome patterns.
- In comparison with the 1 bit History Predictor the Two-bit Saturated Counter will perform better in the long run if the loop is revisited.
  - *The 1-bit history will leave the loop predicting the next execution to be not taken (will be incorrect)*
  - *The two-bit saturated counter will leave the loop predicting the next execution to be taken (will be correct)*

# Example Problem 2

---

For a particular ARM program branch instructions constitute 15% of all the instructions executed, and 60% of them are taken. Assume that the microprocessor has a typical 5 stage pipelined instruction execution unit and a memory latency of 4 cycles. Also assume that wasted cycles are only due control hazards.

- a) Determine the CPI for the above case
- b) Determine the CPI if the branch instruction is substituted with a delayed branch of one delay slot (the instructions following the branch that will always be executed) for which a useful instruction can be placed in the delay slot 75% of the time.
- c) Determine the performance speedup resulting from replacing the original branch instruction with a delayed branch.

# What we know

---

$PB = \% \text{ of instructions that are branches} = 15\% = .15$

$BWC = \text{average number of wasted cycles introduced by branch instructions. Then,}$

$\text{Control}_{\text{PENALTY}} = PB \times BWC$

$CPI = 1 + \text{Control}_{\text{PENALTY}} = 1 + PB \times BWC = 1 + .15 \times BWC$

# Solution of Part a

Since the ARM branch instruction predicts that the branch will not be taken, and 60% of the branches are taken, then, 60% (.6) of the predictions will not be asserted.

The penalty is the aggregate of the cycles wasted due to the 2 instructions cancelled and the 4 idle cycles due to the memory latency (total of 6).

$BWC = (\% \text{ predictions not asserted}) \times (\# \text{ of cycles wasted on a miss prediction})$

$$BWC = .6 \times 6 = 3.6$$

$$CPI = 1 + .15 \times BWC = 1 + .15 \times 3.6 = 1 + .54 = \mathbf{1.54}$$

# Solution of Part b

Since a useful instruction cannot be placed on the delay slot 100% of the time, there will be a penalty due to the delayed branch even if the prediction is asserted.

Thus,  $BWC = \text{Delayed}_{\text{PENALTY}} + \text{Prediction}_{\text{PENALTY}}$

The delayed penalty will be one cycle for the 25% of the branches that cannot have a useful instruction in the delay slot. Then,

$\text{Delayed}_{\text{PENALTY}} = .25 \times 1 = .25 \text{ cycles.}$

This time, every branch that miss the prediction introduces a penalty of 5 cycles (1 instruction canceled and the 4 idle cycles due to the memory latency). Since this happens for 60% of the branches,

$\text{Prediction}_{\text{PENALTY}} = .6 \times 5 = 3 \text{ cycles.}$

Then,  $BWC = .25 + 3 = 3.25 \text{ cycles}$

$CPI = 1 + .15 \times BWC = 1 + .15 \times 3.25 = \mathbf{1.49}$



# Solution of Part c

---

$$\text{Speedup} = \frac{\text{CPI}_{\text{DELAYED BRANCH}} - \text{CPI}_{\text{ARM BRANCH}}}{\text{CPI}_{\text{ARM BRANCH}}}$$

$$\text{Speedup} = (1.49 - 1.54) / 1.54 = - 3.2\%$$

# Lesson Outcomes

- Understand how a pipelined instruction execution unit works
- Know the three types of hazards and how are they detected
- Understand the basic mechanisms to reduce the penalty introduced by hazard interlocks
- Know how to analyze the performance of a processor in terms of CPI